

ps2

April 6, 2024

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os
import seaborn as sns
import tensorflow as tf
from tensorflow.keras import layers, Model
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.applications.resnet50 import preprocess_input
from sklearn.model_selection import train_test_split
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: # Load text data from CSV
df = pd.read_csv('/content/drive/MyDrive/hateful_memes/hateful_memes_original.
↳CSV')

# Adjust image directory path for Colab
image_dir = '/content/drive/MyDrive/hateful_memes/img'
```

```
[ ]: df.head()
```

```
[ ]:
```

	text	img	label
0	a school bus that was engulfed in flames	img/32674.png	0
1	when you ask your dad who is a retired drill s...	img/10246.png	1
2	how i see kim burell everytime she grabs a mic!	img/14570.png	1
3	doing o's with the smoke	img/05316.png	0
4	im gettin white girl wasted tonight	img/20936.png	0

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 10000 entries, 0 to 9999
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    10000 non-null   object
1    img      10000 non-null   object
2    label    10000 non-null   int64
dtypes: int64(1), object(2)
memory usage: 234.5+ KB

```

```
[ ]: df.drop_duplicates(inplace=True)
```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    10000 non-null   object
1    img      10000 non-null   object
2    label    10000 non-null   int64
dtypes: int64(1), object(2)
memory usage: 234.5+ KB

```

```
[ ]: # Preprocess text data
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['text'])
sequences = tokenizer.texts_to_sequences(df['text'])
word_index = tokenizer.word_index
max_sequence_length = max([len(seq) for seq in sequences])
text_data = pad_sequences(sequences, maxlen=max_sequence_length)
```

```
[ ]: text_data
```

```
[ ]: array([[ 0,  0,  0, ..., 5586,  10, 5587],
          [ 0,  0,  0, ..., 3091, 1316,  287],
          [ 0,  0,  0, ..., 3092,   2, 3880],
          ...,
          [ 0,  0,  0, ...,   0,  296,  727],
          [ 0,  0,  0, ...,   80,  363,  333],
          [ 0,  0,  0, ..., 3874,   9,   63]], dtype=int32)
```

```
[ ]: image_data = []
valid_indices = [] # Keep track of valid indices for images
for idx, image_file in enumerate(os.listdir(image_dir)):
    if image_file.endswith('.png'):
```

```

        img_path = os.path.join(image_dir, image_file)
        img = tf.keras.preprocessing.image.load_img(img_path, target_size=(224,
↪224))
        img_array = tf.keras.preprocessing.image.img_to_array(img)
        img_array = preprocess_input(img_array)
        image_data.append(img_array)
        valid_indices.append(idx)
image_data = np.array(image_data)

```

```
[ ]: valid_indices
```

```

[ ]: [0,
      1,
      2,
      3,
      4,
      5,
      6,
      7,
      8,
      9,
      10,
      11,
      12,
      13,
      14,
      15,
      16,
      17,
      18,
      19,
      20,
      21,
      22,
      23,
      24,
      25,
      26,
      27,
      28,
      29,
      30,
      31,
      32,
      33,
      34,
      35,

```

36,
37,
38,
39,
40,
41,
42,
43,
44,
45,
46,
47,
48,
49,
50,
51,
52,
53,
54,
55,
56,
57,
58,
59,
60,
61,
62,
63,
64,
65,
66,
67,
68,
69,
70,
71,
72,
73,
74,
75,
76,
77,
78,
79,
80,
81,
82,

83,
84,
85,
86,
87,
88,
89,
90,
91,
92,
93,
94,
95,
96,
97,
98,
99,
100,
101,
102,
103,
104,
105,
106,
107,
108,
109,
110,
111,
112,
113,
114,
115,
116,
117,
118,
119,
120,
121,
122,
123,
124,
125,
126,
127,
128,
129,

130,
131,
132,
133,
134,
135,
136,
137,
138,
139,
140,
141,
142,
143,
144,
145,
146,
147,
148,
149,
150,
151,
152,
153,
154,
155,
156,
157,
158,
159,
160,
161,
162,
163,
164,
165,
166,
167,
168,
169,
170,
171,
172,
173,
174,
175,
176,

177,
178,
179,
180,
181,
182,
183,
184,
185,
186,
187,
188,
189,
190,
191,
192,
193,
194,
195,
196,
197,
198,
199,
200,
201,
202,
203,
204,
205,
206,
207,
208,
209,
210,
211,
212,
213,
214,
215,
216,
217,
218,
219,
220,
221,
222,
223,

224,
225,
226,
227,
228,
229,
230,
231,
232,
233,
234,
235,
236,
237,
238,
239,
240,
241,
242,
243,
244,
245,
246,
247,
248,
249,
250,
251,
252,
253,
254,
255,
256,
257,
258,
259,
260,
261,
262,
263,
264,
265,
266,
267,
268,
269,
270,

271,
272,
273,
274,
275,
276,
277,
278,
279,
280,
281,
282,
283,
284,
285,
286,
287,
288,
289,
290,
291,
292,
293,
294,
295,
296,
297,
298,
299,
300,
301,
302,
303,
304,
305,
306,
307,
308,
309,
310,
311,
312,
313,
314,
315,
316,
317,

318,
319,
320,
321,
322,
323,
324,
325,
326,
327,
328,
329,
330,
331,
332,
333,
334,
335,
336,
337,
338,
339,
340,
341,
342,
343,
344,
345,
346,
347,
348,
349,
350,
351,
352,
353,
354,
355,
356,
357,
358,
359,
360,
361,
362,
363,
364,

365,
366,
367,
368,
369,
370,
371,
372,
373,
374,
375,
376,
377,
378,
379,
380,
381,
382,
383,
384,
385,
386,
387,
388,
389,
390,
391,
392,
393,
394,
395,
396,
397,
398,
399,
400,
401,
402,
403,
404,
405,
406,
407,
408,
409,
410,
411,

412,
413,
414,
415,
416,
417,
418,
419,
420,
421,
422,
423,
424,
425,
426,
427,
428,
429,
430,
431,
432,
433,
434,
435,
436,
437,
438,
439,
440,
441,
442,
443,
444,
445,
446,
447,
448,
449,
450,
451,
452,
453,
454,
455,
456,
457,
458,

459,
460,
461,
462,
463,
464,
465,
466,
467,
468,
469,
470,
471,
472,
473,
474,
475,
476,
477,
478,
479,
480,
481,
482,
483,
484,
485,
486,
487,
488,
489,
490,
491,
492,
493,
494,
495,
496,
497,
498,
499,
500,
501,
502,
503,
504,
505,

506,
507,
508,
509,
510,
511,
512,
513,
514,
515,
516,
517,
518,
519,
520,
521,
522,
523,
524,
525,
526,
527,
528,
529,
530,
531,
532,
533,
534,
535,
536,
537,
538,
539,
540,
541,
542,
543,
544,
545,
546,
547,
548,
549,
550,
551,
552,

553,
554,
555,
556,
557,
558,
559,
560,
561,
562,
563,
564,
565,
566,
567,
568,
569,
570,
571,
572,
573,
574,
575,
576,
577,
578,
579,
580,
581,
582,
583,
584,
585,
586,
587,
588,
589,
590,
591,
592,
593,
594,
595,
596,
597,
598,
599,

600,
601,
602,
603,
604,
605,
606,
607,
608,
609,
610,
611,
612,
613,
614,
615,
616,
617,
618,
619,
620,
621,
622,
623,
624,
625,
626,
627,
628,
629,
630,
631,
632,
633,
634,
635,
636,
637,
638,
639,
640,
641,
642,
643,
644,
645,
646,

647,
648,
649,
650,
651,
652,
653,
654,
655,
656,
657,
658,
659,
660,
661,
662,
663,
664,
665,
666,
667,
668,
669,
670,
671,
672,
673,
674,
675,
676,
677,
678,
679,
680,
681,
682,
683,
684,
685,
686,
687,
688,
689,
690,
691,
692,
693,

694,
695,
696,
697,
698,
699,
700,
701,
702,
703,
704,
705,
706,
707,
708,
709,
710,
711,
712,
713,
714,
715,
716,
717,
718,
719,
720,
721,
722,
723,
724,
725,
726,
727,
728,
729,
730,
731,
732,
733,
734,
735,
736,
737,
738,
739,
740,

741,
742,
743,
744,
745,
746,
747,
748,
749,
750,
751,
752,
753,
754,
755,
756,
757,
758,
759,
760,
761,
762,
763,
764,
765,
766,
767,
768,
769,
770,
771,
772,
773,
774,
775,
776,
777,
778,
779,
780,
781,
782,
783,
784,
785,
786,
787,

788,
789,
790,
791,
792,
793,
794,
795,
796,
797,
798,
799,
800,
801,
802,
803,
804,
805,
806,
807,
808,
809,
810,
811,
812,
813,
814,
815,
816,
817,
818,
819,
820,
821,
822,
823,
824,
825,
826,
827,
828,
829,
830,
831,
832,
833,
834,

835,
836,
837,
838,
839,
840,
841,
842,
843,
844,
845,
846,
847,
848,
849,
850,
851,
852,
853,
854,
855,
856,
857,
858,
859,
860,
861,
862,
863,
864,
865,
866,
867,
868,
869,
870,
871,
872,
873,
874,
875,
876,
877,
878,
879,
880,
881,

882,
883,
884,
885,
886,
887,
888,
889,
890,
891,
892,
893,
894,
895,
896,
897,
898,
899,
900,
901,
902,
903,
904,
905,
906,
907,
908,
909,
910,
911,
912,
913,
914,
915,
916,
917,
918,
919,
920,
921,
922,
923,
924,
925,
926,
927,
928,

929,
930,
931,
932,
933,
934,
935,
936,
937,
938,
939,
940,
941,
942,
943,
944,
945,
946,
947,
948,
949,
950,
951,
952,
953,
954,
955,
956,
957,
958,
959,
960,
961,
962,
963,
964,
965,
966,
967,
968,
969,
970,
971,
972,
973,
974,
975,

```
976,  
977,  
978,  
979,  
980,  
981,  
982,  
983,  
984,  
985,  
986,  
987,  
988,  
989,  
990,  
991,  
992,  
993,  
994,  
995,  
996,  
997,  
998,  
999,  
...]
```

```
[ ]: text_data_filtered = text_data[valid_indices]  
y_filtered = df['label'].iloc[valid_indices]
```

```
[ ]: X_text = text_data_filtered  
X_image = image_data  
y = y_filtered.values
```

```
[ ]: X_text_train, X_text_test, X_image_train, X_image_test, y_train, y_test =  
    ↪train_test_split(  
        X_text, X_image, y, test_size=0.2, random_state=42)  
X_text_train, X_text_val, X_image_train, X_image_val, y_train, y_val =  
    ↪train_test_split(  
        X_text_train, X_image_train, y_train, test_size=0.1, random_state=42)
```

```
[ ]: text_input = layers.Input(shape=(max_sequence_length,), name='text_input')  
text_embedding = layers.Embedding(len(word_index) + 1, 100,  
    ↪input_length=max_sequence_length)(text_input)  
text_flatten = layers.Flatten()(text_embedding)  
  
image_input = layers.Input(shape=(224, 224, 3), name='image_input')
```



```

base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224,
↳224, 3))
image_output = base_model(image_input)
image_flatten = layers.Flatten()(image_output)

concatenated = layers.concatenate([text_flatten, image_flatten])
output = layers.Dense(1, activation='sigmoid')(concatenated)

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 [=====] - 1s 0us/step

```
[ ]: model = Model(inputs=[text_input, image_input], outputs=output)
```

```
[ ]: model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])
```

```
[ ]: model.fit([X_text_train, X_image_train], y_train, validation_data=([X_text_val,
↳X_image_val], y_val), epochs=10, batch_size=32)
```

Epoch 1/10

69/69 [=====] - 1865s 27s/step - loss: 2.3896 - accuracy: 0.5852 - val_loss: 5406.7153 - val_accuracy: 0.6122

Epoch 2/10

69/69 [=====] - 1834s 27s/step - loss: 1.0522 - accuracy: 0.6280 - val_loss: 22.0468 - val_accuracy: 0.6082

Epoch 3/10

69/69 [=====] - 1784s 26s/step - loss: 0.6331 - accuracy: 0.7391 - val_loss: 2.1582 - val_accuracy: 0.5265

Epoch 4/10

69/69 [=====] - 1715s 25s/step - loss: 0.5437 - accuracy: 0.8506 - val_loss: 0.8670 - val_accuracy: 0.5633

Epoch 5/10

69/69 [=====] - 1694s 25s/step - loss: 0.3645 - accuracy: 0.9030 - val_loss: 2.1191 - val_accuracy: 0.4816

Epoch 6/10

69/69 [=====] - 1695s 25s/step - loss: 0.2139 - accuracy: 0.9317 - val_loss: 0.8870 - val_accuracy: 0.5959

Epoch 7/10

69/69 [=====] - 1674s 24s/step - loss: 0.1662 - accuracy: 0.9545 - val_loss: 1.1787 - val_accuracy: 0.4980

Epoch 8/10

69/69 [=====] - 1712s 25s/step - loss: 0.1234 - accuracy: 0.9604 - val_loss: 1.0449 - val_accuracy: 0.5918

Epoch 9/10

69/69 [=====] - 1686s 24s/step - loss: 0.0557 - accuracy: 0.9891 - val_loss: 1.0527 - val_accuracy: 0.5878

```
[ ]: <keras.src.callbacks.History at 0x7c016cb8e980>
```

[illegible]

[1],
[1],
[1],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[1],
[1],
[1],
[1],
[0],
[0],
[0],
[1]

[1],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[0],
[1],
[0],
[0],
[0],
[1],
[1],
[1],
[0],
[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[1],
[1],
[1],
[1],
[1],
[0],
[0],
[1],

[illegible]

[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[1],
[1],
[1],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[1],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0]

[1],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0]

[0],
[0],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[1],
[1],
[1],
[0],
[1],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[1],
[1],
[1],

[1],
[0],
[1],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],

[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[1],
[1],
[1],
[0],
[1],
[1],
[1],
[1],
[1],
[1],
[1]

[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[1],
[1],
[1],
[0],
[0],
[1],
[1],
[1],
[0],
[0],
[1],
[1],
[1],
[1],
[0],
[0],
[1],
[1],
[0],
[0],
[1],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[0]

[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[1],
[1],
[0],
[1],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],

[0],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[1],
[1],
[1],
[1],
[1],
[1],
[0],
[1],
[0],
[0],
[0],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[0],
[0],
[1],

[illegible]

```
[1],
[1],
[0],
[0],
[1],
[0],
[0],
[0],
[1],
[1],
[0],
[0],
[0],
[1],
[0],
[1],
[0],
[0],
[0],
[0]])
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
      f1_score, roc_auc_score, confusion_matrix
```

```
[ ]: precision = precision_score(y_test, y_pred)
      recall = recall_score(y_test, y_pred)
      f1 = f1_score(y_test, y_pred)
      auc = roc_auc_score(y_test, y_pred_prob)

      print(f'Precision: {precision}, Recall: {recall}, F1-score: {f1}, AUC: {auc}')
```

```
Precision: 0.4973821989528796, Recall: 0.4460093896713615, F1-score:
0.4702970297029703, AUC: 0.6200721919456437
```

```
[ ]: conf_matrix = confusion_matrix(y_test, y_pred)
```

```
[ ]: print(conf_matrix)
```

```
[[302  96]
 [118  95]]
```

```
[ ]: plt.figure(figsize=(8, 6))
      sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False,
                  xticklabels=['Non-harmful', 'Harmful'], yticklabels=['Non-harmful',
                  'Harmful'])
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')  
plt.show()
```

